

jupyter-train2

March 28, 2023

```
[4]: #!/usr/bin/python3
# -*- coding: utf-8 -*-
# @Time      : 2023/3/27 21:48
# @Author    : Genius_limeng
# @FileName  : train.py
# @Software  : PyCharm

import random
import os
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.layers import Convolution2D, MaxPooling2D, AveragePooling2D
from keras.optimizers import SGD,adam
from keras.utils import np_utils
from keras.models import load_model
from keras import backend as K
from load_dataset import load_dataset, resize_image
from keras.models import Model
from sklearn.metrics import classification_report
import numpy as np
import seaborn as sns
from sklearn.metrics import confusion_matrix

IMAGE_SIZE = 128

#
class Dataset:
    def __init__(self, path_name):
        #
        self.train_images = None
        self.train_labels = None
        #
        self.valid_images = None
        self.valid_labels = None
        #
```

```

self.test_images = None
self.test_labels = None
#
self.path_name = path_name
#
self.user_num = len(os.listdir(path_name))
#
self.input_shape = None

#
def load(self, img_rows=IMAGE_SIZE, img_cols=IMAGE_SIZE, img_channels=3):
    #
    nb_classes = self.user_num
    #
    images, labels = load_dataset(self.path_name)

    train_images, valid_images, train_labels, valid_labels =
↪train_test_split(images, labels, test_size=0.3,
                                                            random_state=random.
↪randint(0, 100))
    _, test_images, _, test_labels = train_test_split(images, labels,
↪test_size=0.3,
                                                            random_state=random.
↪randint(0, 100))

    #      'th'      channels,rows,cols :rows,cols,channels
    #      keras
    if K.image_dim_ordering() == 'th':
        train_images = train_images.reshape(train_images.shape[0],
↪img_channels, img_rows, img_cols)
        valid_images = valid_images.reshape(valid_images.shape[0],
↪img_channels, img_rows, img_cols)
        test_images = test_images.reshape(test_images.shape[0],
↪img_channels, img_rows, img_cols)
        self.input_shape = (img_channels, img_rows, img_cols)
    else:
        train_images = train_images.reshape(train_images.shape[0],
↪img_rows, img_cols, img_channels)
        valid_images = valid_images.reshape(valid_images.shape[0],
↪img_rows, img_cols, img_channels)
        test_images = test_images.reshape(test_images.shape[0], img_rows,
↪img_cols, img_channels)
        self.input_shape = (img_rows, img_cols, img_channels)

    #
    print(train_images.shape[0], 'train samples')

```

```

print(valid_images.shape[0], 'valid samples')
print(test_images.shape[0], 'test samples')

#      categorical_crossentropy      nb_classes
#      one-hot      4
train_labels = np_utils.to_categorical(train_labels, nb_classes)
valid_labels = np_utils.to_categorical(valid_labels, nb_classes)
test_labels = np_utils.to_categorical(test_labels, nb_classes)

#
train_images = train_images.astype('float32')
valid_images = valid_images.astype('float32')
test_images = test_images.astype('float32')

#      ,      0~1
train_images /= 255
valid_images /= 255
test_images /= 255

self.train_images = train_images
self.valid_images = valid_images
self.test_images = test_images
self.train_labels = train_labels
self.valid_labels = valid_labels
self.test_labels = test_labels

# CNN
class Model:
    def __init__(self):
        self.model = None
        self.history = None

    #
    def build_model(self, dataset, nb_classes=4):
        #
        self.model = Sequential()
        #      CNN      add
        # -----1-----
        self.model.add(Convolution2D(32, 3, 3, border_mode='same',
                                     input_shape=dataset.input_shape)) # 32
        self.model.add(Convolution2D(32, 3, 3, border_mode='same')) # 32
        self.model.add(Activation('relu'))
        self.model.add(MaxPooling2D(pool_size=(2, 2))) #
        self.model.add(Dropout(0.25)) # Dropout
        # -----2-----

```

```

self.model.add(Convolution2D(64, 3, 3, border_mode='same')) # 64
self.model.add(Convolution2D(64, 3, 3, border_mode='same')) # 64
self.model.add(Activation('relu'))
self.model.add(MaxPooling2D(pool_size=(2, 2))) #
self.model.add(Dropout(0.25)) # Dropout
# -----3-----
self.model.add(Convolution2D(128, 3, 3, border_mode='same')) # 128
self.model.add(Convolution2D(128, 3, 3, border_mode='same')) # 128
self.model.add(Convolution2D(128, 3, 3, border_mode='same')) # 128
self.model.add(Activation('relu'))
self.model.add(MaxPooling2D(pool_size=(2, 2))) #
self.model.add(Dropout(0.25)) # Dropout
# -----4-----
self.model.add(Convolution2D(256, 3, 3, border_mode='same')) # 256
self.model.add(Convolution2D(256, 3, 3, border_mode='same')) # 256
self.model.add(Convolution2D(256, 3, 3, border_mode='same')) # 256
self.model.add(Activation('relu'))
self.model.add(MaxPooling2D(pool_size=(2, 2))) #
self.model.add(Dropout(0.25)) # Dropout
# -----5-----
self.model.add(Convolution2D(256, 3, 3, border_mode='same')) # 256
self.model.add(Convolution2D(256, 3, 3, border_mode='same')) # 256
self.model.add(Convolution2D(256, 3, 3, border_mode='same')) # 256
self.model.add(Activation('relu'))
self.model.add(MaxPooling2D(pool_size=(2, 2))) #
# self.model.add(AveragePooling2D(pool_size=(2, 2), strides=2)) #
self.model.add(Dropout(0.25)) # Dropout
# -----6-----
self.model.add(Flatten()) # Flatten
self.model.add(Dense(256)) # Dense ,
self.model.add(Activation('relu'))
self.model.add(Dropout(0.5)) # Dropout
self.model.add(Dense(nb_classes)) # Dense
self.model.add(Activation('softmax')) #

#
self.model.summary()

#
def compile(self):
    # SGD (learning_rate)
    # 0.01 (decay)
    # (momentum),
    #
    #
    sgd = SGD(lr=0.0001, decay=1e-6, # 10 -6
              momentum=0.8, nesterov=True) # SGD+momentum

```

```

Adam = adam(lr=0.0001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.
↪0, amsgrad=False)
self.model.compile(loss='categorical_crossentropy',
                    optimizer=sgd,
                    metrics=['accuracy']) #

#
def train(self, dataset, batch_size=20, nb_epoch=20, ↵
↪data_augmentation=True):

    if not data_augmentation:
        self.history = self.model.fit(dataset.train_images,
                                       dataset.train_labels,
                                       batch_size=batch_size,
                                       nb_epoch=nb_epoch,
                                       validation_data=(dataset.
↪valid_images, dataset.valid_labels),
                                       shuffle=True)

    #
    else:
        #          datagen datagen
        #          python
        datagen = ImageDataGenerator(
            featurewise_center=False, #          0
            samplewise_center=False, #          0
            featurewise_std_normalization=False, #
            samplewise_std_normalization=False, #
            zca_whitening=False, #          ZCA
            rotation_range=20, #          ( 0 180)
            width_shift_range=0.2, #          0~1
            height_shift_range=0.2, #
            horizontal_flip=True, #
            vertical_flip=False) #

        #          ZCA
        datagen.fit(dataset.train_images)

        #
        self.history = self.model.fit_generator(datagen.flow(dataset.
↪train_images, dataset.train_labels,
                                                                    ↵
↪batch_size=batch_size),
                                                samples_per_epoch=dataset.
↪train_images.shape[0],
                                                nb_epoch=nb_epoch,

```

```

validation_data=(dataset.
↪valid_images, dataset.valid_labels))

MODEL_PATH = './model/CNN-FaceRecognitionModel.h5'

#
def save_model(self, file_path=MODEL_PATH):
    self.model.save(file_path)

#
def load_model(self, file_path=MODEL_PATH):
    self.model = load_model(file_path)

#
def evaluate(self, dataset):
    score = self.model.evaluate(dataset.test_images, dataset.test_labels,
↪batch_size=20, verbose=1)
    print("%s: %.2f%%" % (self.model.metrics_names[1], score[1] * 100))
    print("score:", score)
    plt.plot(self.history.history['acc'])
    plt.plot(self.history.history['val_acc'])
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'val'], loc='upper left')
    plt.show()

    plt.plot(self.history.history['loss'])
    plt.plot(self.history.history['val_loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'val'], loc='upper left')
    plt.show()

#
def predict(self, dataset):
    prob = self.model.predict(dataset.test_images)
    predIdxs = np.argmax(prob, axis=1)
    print('\n')
    print(classification_report(dataset.test_labels.argmax(axis=1),
↪predIdxs,
                                =['faces_other', 'LiMeng', 'ShiWei',
↪'YangXinYe', 'ZhengLiWei'], digits=5))

#
def heatmap(self, dataset):

```

```

prob = self.model.predict(dataset.test_images)
predIdxs = np.argmax(prob, axis=1)
cm = confusion_matrix(dataset.test_labels.argmax(axis=1), predIdxs)
sns.heatmap(cm, annot=True)
plt.title('Confusion matrix')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

#
def face_predict(self, image):
    #
    if K.image_dim_ordering() == 'th' and image.shape != (1, 3, IMAGE_SIZE, IMAGE_SIZE):
        image = resize_image(image) # IMAGE_SIZE x IMAGE_SIZE
        image = image.reshape((1, 3, IMAGE_SIZE, IMAGE_SIZE)) #
    elif K.image_dim_ordering() == 'tf' and image.shape != (1, IMAGE_SIZE, IMAGE_SIZE, 3):
        image = resize_image(image)
        image = image.reshape((1, IMAGE_SIZE, IMAGE_SIZE, 3))

    #
    image = image.astype('float32')
    image /= 255

    #
    result_probability = self.model.predict_proba(image)
    print('result:', result_probability)

    #
    if max(result_probability[0]) >= 0.8:
        result = self.model.predict_classes(image)
        print('result:', result)
        #
        return result[0]
    else:
        print('result:none')
        return -1

```

```

[5]: if __name__ == '__main__':
    user_num = len(os.listdir('./dataset/'))
    # print(" ", user_num, " ")
    dataset = Dataset('./dataset/')
    dataset.load()
    cnn = Model()
    cnn.build_model(dataset, nb_classes=user_num)
    cnn.compile() #

```

```
#
cnn.train(dataset)
#
cnn.save_model(file_path='./model/CNN-FaceRecognitionModel.h5')
```

```
36373 train samples
15589 valid samples
15589 test samples
```

```
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:115:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(32, (3, 3),
input_shape=(128, 128,..., padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:116:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(32, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:121:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(64, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:122:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(64, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:127:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(128, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:128:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(128, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:129:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(128, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:134:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(256, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:135:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(256, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:136:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(256, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:141:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(256, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:142:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(256, (3, 3),
padding="same")`
C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:143:
UserWarning: Update your `Conv2D` call to the Keras 2 API: `Conv2D(256, (3, 3),
padding="same")`
```


Layer (type)	Output Shape	Param #
conv2d_14 (Conv2D)	(None, 128, 128, 32)	896
conv2d_15 (Conv2D)	(None, 128, 128, 32)	9248
activation_2 (Activation)	(None, 128, 128, 32)	0
max_pooling2d_6 (MaxPooling2D)	(None, 64, 64, 32)	0
dropout_7 (Dropout)	(None, 64, 64, 32)	0
conv2d_16 (Conv2D)	(None, 64, 64, 64)	18496
conv2d_17 (Conv2D)	(None, 64, 64, 64)	36928
activation_3 (Activation)	(None, 64, 64, 64)	0
max_pooling2d_7 (MaxPooling2D)	(None, 32, 32, 64)	0
dropout_8 (Dropout)	(None, 32, 32, 64)	0
conv2d_18 (Conv2D)	(None, 32, 32, 128)	73856
conv2d_19 (Conv2D)	(None, 32, 32, 128)	147584
conv2d_20 (Conv2D)	(None, 32, 32, 128)	147584
activation_4 (Activation)	(None, 32, 32, 128)	0
max_pooling2d_8 (MaxPooling2D)	(None, 16, 16, 128)	0
dropout_9 (Dropout)	(None, 16, 16, 128)	0
conv2d_21 (Conv2D)	(None, 16, 16, 256)	295168
conv2d_22 (Conv2D)	(None, 16, 16, 256)	590080
conv2d_23 (Conv2D)	(None, 16, 16, 256)	590080
activation_5 (Activation)	(None, 16, 16, 256)	0
max_pooling2d_9 (MaxPooling2D)	(None, 8, 8, 256)	0
dropout_10 (Dropout)	(None, 8, 8, 256)	0
conv2d_24 (Conv2D)	(None, 8, 8, 256)	590080

conv2d_25 (Conv2D)	(None, 8, 8, 256)	590080
conv2d_26 (Conv2D)	(None, 8, 8, 256)	590080
activation_6 (Activation)	(None, 8, 8, 256)	0
max_pooling2d_10 (MaxPooling)	(None, 4, 4, 256)	0
dropout_11 (Dropout)	(None, 4, 4, 256)	0
flatten_2 (Flatten)	(None, 4096)	0
dense_3 (Dense)	(None, 256)	1048832
activation_7 (Activation)	(None, 256)	0
dropout_12 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 5)	1285
activation_8 (Activation)	(None, 5)	0

=====
Total params: 4,730,277

Trainable params: 4,730,277

Non-trainable params: 0

C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:207:
UserWarning: The semantics of the Keras 2 argument `steps_per_epoch` is not the same as the Keras 1 argument `samples_per_epoch`. `steps_per_epoch` is the number of batches to draw from the generator at each epoch. Basically
steps_per_epoch = samples_per_epoch/batch_size. Similarly
`nb_val_samples`->`validation_steps` and `val_samples`->`steps` arguments have changed. Update your method calls accordingly.

C:\Users\limeng\anaconda3\envs\py37\lib\site-packages\ipykernel_launcher.py:207:
UserWarning: Update your `fit_generator` call to the Keras 2 API:
`fit_generator(<keras_pre..., validation_data=(array([[...],
steps_per_epoch=1818, epochs=20))`

Epoch 1/20

1818/1818 [=====] - 164s 90ms/step - loss: 1.2440 -
acc: 0.5390 - val_loss: 1.2831 - val_acc: 0.6378

Epoch 2/20

1818/1818 [=====] - 152s 84ms/step - loss: 0.9507 -
acc: 0.6319 - val_loss: 0.9635 - val_acc: 0.6784

Epoch 3/20

1818/1818 [=====] - 150s 83ms/step - loss: 0.7966 -
acc: 0.6872 - val_loss: 0.7528 - val_acc: 0.7732

Epoch 4/20
1818/1818 [=====] - 160s 88ms/step - loss: 0.7087 -
acc: 0.7277 - val_loss: 0.8690 - val_acc: 0.6642

Epoch 5/20
1818/1818 [=====] - 155s 85ms/step - loss: 0.6311 -
acc: 0.7644 - val_loss: 0.5631 - val_acc: 0.8323

Epoch 6/20
1818/1818 [=====] - 153s 84ms/step - loss: 0.5546 -
acc: 0.7996 - val_loss: 0.5457 - val_acc: 0.8058

Epoch 7/20
1818/1818 [=====] - 152s 84ms/step - loss: 0.4817 -
acc: 0.8270 - val_loss: 0.3507 - val_acc: 0.8905

Epoch 8/20
1818/1818 [=====] - 152s 84ms/step - loss: 0.4285 -
acc: 0.8476 - val_loss: 0.3360 - val_acc: 0.8881

Epoch 9/20
1818/1818 [=====] - 151s 83ms/step - loss: 0.3762 -
acc: 0.8663 - val_loss: 0.2464 - val_acc: 0.9234

Epoch 10/20
1818/1818 [=====] - 148s 82ms/step - loss: 0.3331 -
acc: 0.8847 - val_loss: 0.2153 - val_acc: 0.9299

Epoch 11/20
1818/1818 [=====] - 148s 82ms/step - loss: 0.2930 -
acc: 0.8984 - val_loss: 0.1871 - val_acc: 0.9397

Epoch 12/20
1818/1818 [=====] - 152s 84ms/step - loss: 0.2639 -
acc: 0.9091 - val_loss: 0.1430 - val_acc: 0.9585

Epoch 13/20
1818/1818 [=====] - 152s 84ms/step - loss: 0.2384 -
acc: 0.9182 - val_loss: 0.1227 - val_acc: 0.9658

Epoch 14/20
1818/1818 [=====] - 152s 84ms/step - loss: 0.2121 -
acc: 0.9271 - val_loss: 0.1113 - val_acc: 0.9695

Epoch 15/20
1818/1818 [=====] - 155s 85ms/step - loss: 0.1922 -
acc: 0.9340 - val_loss: 0.0900 - val_acc: 0.9749

Epoch 16/20
1818/1818 [=====] - 157s 86ms/step - loss: 0.1812 -
acc: 0.9409 - val_loss: 0.0779 - val_acc: 0.9805

Epoch 17/20
1818/1818 [=====] - 173s 95ms/step - loss: 0.1586 -
acc: 0.9486 - val_loss: 0.1092 - val_acc: 0.9677


Epoch 18/20
1818/1818 [=====] - 155s 85ms/step - loss: 0.1505 -
acc: 0.9509 - val_loss: 0.0582 - val_acc: 0.9858


Epoch 19/20
1818/1818 [=====] - 155s 85ms/step - loss: 0.1418 -
acc: 0.9537 - val_loss: 0.0547 - val_acc: 0.9877

Epoch 20/20
1818/1818 [=====] - 174s 96ms/step - loss: 0.1347 -
acc: 0.9581 - val_loss: 0.0499 - val_acc: 0.9872

```
[11]: #  
      cnn.evaluate(dataset)
```

15589/15589 [=====] - 18s 1ms/step
acc: 98.73%
score: [0.04902047472949644, 0.9872987334131315]






```
[9]: #
      cnn.predict(dataset)
```

	precision	recall	f1-score	support
LiMeng	0.99976	0.99030	0.99501	8455
ShiWei	0.94357	0.99945	0.97071	1807
YangXinYe	0.97567	0.96064	0.96809	1753
ZhengLiWei	0.97920	0.98459	0.98189	1817
faces_other	0.99542	0.98976	0.99258	1757
accuracy			0.98730	15589
macro avg	0.97872	0.98495	0.98166	15589
weighted avg	0.98765	0.98730	0.98736	15589

```
[10]: #
       cnn.heatmap(dataset)
```



[]: